

Test 3: Performance between Indexing and Pointer Arithmetic In  
clearing Array and Performance of Dot product between Indexing,  
Pointers and Vector Instructions.

By Demetrios Doumas  
Fall 2016  
CSC 343  
Prof. Gertner

## Table of Contents

1. Objective.....	pg 3
2. Indexing VS Pointers in Visual Studio 2013 Platform.....	pgs 3-13
3. Indexing VS Pointers in Linux Platform.....	pgs 14-19
4. Dot Product Indexing VS Pointers in Visual Studio.....	pgs 20-26
5. Dot Product Computation Using Vector Instruction.....	pgs 27-29
6. Conclusion.....	pg 30
7. Appendix	

**Objective:** Compare performance of clearing an array and dot product by indexing arithmetic and by clearing the array using pointer arithmetic. Also, compare the performance of a dot product scalar code computation with the vector instruction code.

## **Indexing VS Pointers in Visual Studio 2013 Platform:**

Using Indexing Arithmetic:

Main.cpp file:

```
01. #include <windows.h>
02. #include<iostream>
03. using namespace std;
04. //using namespace System;
05.
06. //void main(int[], int);
07. void ClearUsingIndex(int Array[], int size);
08.
09. static int Array[10] = { 1, 2, 3, 4, 5, 6, 7, 8, 9, -1 };
10.
11. int main(){
12.     int size = 10;
13.     ClearUsingIndex(Array, size);
14.
15.     return 0;
16. }
```

IndexArrayClear.cpp file:

```
01. void ClearUsingIndex(int Array[], int size) {
02.     int i;
03.     for (i = 0; i < size; i += 1)
04.         Array[i] = 0;
05. }
```

## Compiler generated Assembly Code listing of Clear Array using Index function:

```

01. ; Listing generated by Microsoft (R) Optimizing Compiler Version 18.00.21005.1
02.
03. TITLE c:\Users\Demetri\documents\visual studio 2013\Projects\IndexArrayClear\IndexArrayClear\indexclear.cpp
04. .686P
05. .XMM
06. include listing.inc
07. .model flat
08.
09. INCLUDELIB MSVCRTD
10. INCLUDELIB OLDNAMES
11.
12. PUBLIC ?ClearUsingIndex@@YAXQAHH@Z ; ClearUsingIndex
13. EXTRN __RTC_InitBase:PROC
14. EXTRN __RTC_Shutdown:PROC
15. ; COMDAT rtc$TMZ
16. rtc$TMZ SEGMENT
17. ;__RTC_Shutdown.rtc$TMZ DD FLAT:__RTC_Shutdown
18. rtc$TMZ ENDS
19. ; COMDAT rtc$IMZ
20. rtc$IMZ SEGMENT
21. ;__RTC_InitBase.rtc$IMZ DD FLAT:__RTC_InitBase
22. rtc$IMZ ENDS
23. ; Function compile flags: /Odtp /RTCsu /ZI
24. ; COMDAT ?ClearUsingIndex@@YAXQAHH@Z
25. _TEXT SEGMENT
26. _i$ = -8 ; size = 4
27. _Array$ = 8 ; size = 4
28. _size$ = 12 ; size = 4
29. ?ClearUsingIndex@@YAXQAHH@Z PROC ; ClearUsingIndex, COMDAT
30. ; File c:\users\demetri\documents\visual studio 2013\projects\indexarrayclear\indexarrayclear\indexclear.cpp
31. ; line 1
32. push ebp
33. mov ebp, esp
34. sub esp, 204 ; 000000cch
35. push ebx
36. push esi
37. push edi
38. lea edi, DWORD PTR [ebp-204]
39. mov ecx, 51 ; 00000033H
40. mov eax, -858993460 ; ccccccccH

```

continue...

```

40. mov eax, -858993460 ; ccccccccH
41. rep stosd
42. ; line 3
43. mov DWORD PTR _i$[ebp], 0
44. jmp SHORT $LN3@ClearUsing
45. $LN2@ClearUsing:
46. mov eax, DWORD PTR _i$[ebp]
47. add eax, 1
48. mov DWORD PTR _i$[ebp], eax
49. $LN3@ClearUsing:
50. mov eax, DWORD PTR _i$[ebp]
51. cmp eax, DWORD PTR _size$[ebp]
52. jge SHORT $LN4@ClearUsing
53. ; line 4
54. mov eax, DWORD PTR _i$[ebp]
55. mov ecx, DWORD PTR _Array$[ebp]
56. mov DWORD PTR [ecx+eax*4], 0
57. jmp SHORT $LN2@ClearUsing
58. $LN4@ClearUsing:
59. ; line 5
60. pop edi
61. pop esi
62. pop ebx
63. mov esp, ebp
64. pop ebp
65. ret 0
66. ?ClearUsingIndex@@YAXQAHH@Z ENDP ; ClearUsingIndex
67. _TEXT ENDS
68. END

```

Functionality of Compiler generated Assembly Code listing of Clear Array using Index in the last page:

In line 26 to 28, the compiler stores the offset of the index of the loop, memory address of Array, and the size array at those particular offset from the base pointer of the stack. In line 32 to 41 the stack is being implemented. The index value in the for-loop gets initialized to one in line 43. The condition of the for-loop is translated in lines 50 to 52. The value at `_i` is the offset from the base pointer that store the index value `I` in the for-loop is store to register `eax`. Then it gets compared with the size of the array at an offset of `_size` from the base pointer. In line 4, the value of `I` in the for loop gets stored to `eax` and the beginning memory address is stored to the `ecx` register by the use of the offset `_Array` from the base pointer. The memory address in `ecx` get incremented by `I` times by 4 bytes, and initialize it to zero. The pattern then repeats until in line 3 where the compare instruction compares the index value at `eax` to the size of the array and if it is greater than then the instruction pointer then starts to deallocated the stack frame and bring the function to the end.

Optimized Assembly Code of Clear Array using Index function:

```

01. .686P
02. .XMM
03. .model flat
04.
05. _TEXT SEGMENT
06. _i$ = -8 ; size = 4
07. _Array$ = 8 ; size = 4
08. _size$ = 12 ; size = 4
09. ?ClearUsingIndex@@YAXQAHH@Z PROC ; ClearUsingIndex, COMDAT
10.
11. ; Line 1
12. push ebp
13. mov ebp, esp
14. sub esp, 204 ; 000000cch
15. push ebx
16. push esi
17. push edi
18. lea edi, DWORD PTR [ebp-204]
19. mov ecx, 51 ; 00000033H
20. mov eax, -858993460 ; ccccccccH
21. rep stosd
22.
23. ;9 instructions
24. ;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;
25. mov DWORD PTR _i$[ebp], 0
26. mov eax, DWORD PTR _i$[ebp]
27. $LN2@ClearUsing:
28. cmp eax, DWORD PTR _size$[ebp]
29. jge short exit
30. mov eax, DWORD PTR _i$[ebp]
31. mov ecx, DWORD PTR _Array$[ebp]
32. mov DWORD PTR [ecx + eax*4], 0
33. add DWORD PTR _i$[ebp], 1
34. jmp SHORT $LN2@ClearUsing
35. ;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;
36. exit:
37.
38.

```

Continued...

```
36.
37.     exit:
38.
39. ;Line 5
40.     pop edi
41.     pop esi
42.     pop ebx
43.     mov esp, ebp
44.     pop ebp
45.     ret 0
46. ?ClearUsingIndex@@YAXQAHH@Z ENDP          ; ClearUsingIndex
47. _TEXT    ENDS
48. END
```

The code above is the optimized version from the compiled generated code. The optimized code has 9 lines of code for the body, not including the stack code, and the compilers version takes 12 lines. The first part of the code is to initialize the index *i* to zero. Move index (*i*) to the *eax* register. Then compare *eax* with the size of the array length. If *eax* is greater than the size of the array then jump to the remove stack code towards the end of the function. In the body of the loop we are instructing the compiler to generate array *A* at a certain index then initialize it to zero (*A*[*i*]=0). This happens by getting access to the current index and the address first (lines 30 and 32). Line 33 increments *i* by one. Then the procedure jumps back to loop and the body of the assemble code is repeated until *eax* is above the size of the Array being passed to the function.

Using Pointer Arithmetic:

Main.cpp

```

01. #include<iostream>
02. using namespace std;
03.
04. void ClearUsingPointers(int *array, int size);
05.
06. static int Array[10] = { 1, 2, 3, 4, 5, 6, 7, 8, 9, -1 };
07.
08. int main()
09. {
10.     int size = 10;
11.
12.     for (int i = 0; i < 10; i++)
13.     {
14.         cout << Array[i] << endl;
15.     }
16.
17.     cout << endl;
18.     cout << endl;
19.     ClearUsingPointers(Array, size);
20.
21.
22.     for (int i = 0; i < 10; i++)
23.     {
24.         cout << Array[i] << endl;
25.     }
26.
27.
28.     return 0;
29. }
30.

```

PointerArrayClear.cpp file:

```

01. void ClearUsingPointers(int *array, int size) {
02.     int *p;
03.     for (p = &array[0]; p < &array[size]; p = p + 1)
04.         *p = 0;
05. }

```

Compiler generated Assembly Code listing of Clear Array using Pointer function:

```

01. ; Listing generated by Microsoft (R) Optimizing Compiler Version 18.00.21005.1
02.
03.     TITLE    c:\Users\Demetri\documents\visual studio 2013\Projects\PointerArrayClear\PointerArrayClear\PointerArrayClear.cpp
04.     .686P
05.     .XMM
06.     include listing.inc
07.     .model    flat
08.
09. INCLUDELIB MSVCRTD
10. INCLUDELIB OLDNAMES
11.
12. PUBLIC  ?ClearUsingPointers@@YAXPAHH@Z      ; ClearUsingPointers
13. EXTRN  __RTC_InitBase:PROC
14. EXTRN  __RTC_Shutdown:PROC
15. ; COMDAT rtc$TMZ
16. rtc$TMZ SEGMENT
17. ;__RTC_Shutdown.rtc$TMZ DD FLAT:__RTC_Shutdown
18. rtc$TMZ ENDS
19. ; COMDAT rtc$IMZ
20. rtc$IMZ SEGMENT
21. ;__RTC_InitBase.rtc$IMZ DD FLAT:__RTC_InitBase
22. rtc$IMZ ENDS
23. ; Function compile flags: /Odtp /RTCsu /ZI
24. ; COMDAT ?ClearUsingPointers@@YAXPAHH@Z
25. _TEXT SEGMENT
26. _p$ = -8 ; size = 4
27. _array$ = 8 ; size = 4
28. _size$ = 12 ; size = 4
29. ?ClearUsingPointers@@YAXPAHH@Z PROC ; ClearUsingPointers, COMDAT
30. ; File c:\users\demetri\documents\visual studio 2013\projects\pointerarrayclear\pointerarrayclear\pointerarrayclear.cpp
31. ; Line 1
32.     push    ebp
33.     mov     ebp, esp
34.     sub     esp, 204 ; 000000cch
35.     push    ebx
36.     push    esi
37.     push    edi
38.     lea     edi, DWORD PTR [ebp-204]
39.     mov     ecx, 51 ; 00000033H
40.     mov     eax, -858993460 ; ccccccccH

```

Continued...

```

41.     rep stosd
42. ; Line 3
43.     mov eax, 4
44.     imul ecx, eax, 0
45.     add ecx, DWORD PTR _array$[ebp]
46.     mov DWORD PTR _p$[ebp], ecx
47.     jmp SHORT $LN3@ClearUsing
48. $LN2@ClearUsing:
49.     mov eax, DWORD PTR _p$[ebp]
50.     add eax, 4
51.     mov DWORD PTR _p$[ebp], eax
52. $LN3@ClearUsing:
53.     mov eax, DWORD PTR _size$[ebp]
54.     mov ecx, DWORD PTR _array$[ebp]
55.     lea edx, DWORD PTR [ecx+eax*4]
56.     cmp DWORD PTR _p$[ebp], edx
57.     jae SHORT $LN4@ClearUsing
58. ; Line 4
59.     mov eax, DWORD PTR _p$[ebp]
60.     mov DWORD PTR [eax], 0
61.     jmp SHORT $LN2@ClearUsing
62. $LN4@ClearUsing:
63. ; Line 5
64.     pop edi
65.     pop esi
66.     pop ebx
67.     mov esp, ebp
68.     pop ebp
69.     ret 0
70. ?ClearUsingPointers@@YAXPAHH@Z ENDP ; ClearUsingPointers
71. _TEXT ENDS
72. END

```

In the last page is the visual studio compiler generated assembly code. The first line is the creation of the stack. The third line is where the body of the function is executed. There were a few step involved at line 3. The value 4 gets copied to the register eax. The register ecx is initialized with a value of zero. The next two lines, the register get the first address of the array and store it to the compiler variable \_p\$. A jump is executed to \$LN3@ClearUsing. The array size is stored to register eax. The memory address of array size is stored in ecx register. The last address is computed and loaded to register edx. The address of \_p\$ is compared with the last address of the array. If the current address is not the same then the next line after the jump instruction is executed. In line 4, the value in \_p\$ is stored at eax. Then zero is moved to eax. The next instruction jumps to \$LN2@ClearUsing. This is where the address is move to the next index by moving the value of \_p\$ to eax and then add 4. Then the address in eax gets saved back to \_p\$. This pattern repeats until eax is equal to edx or the last address of the array. If they are equal then the jump to remove the stack frame is called and the function ends.



## Optimized generated Assembly compiler Code listing of Clear Array using pointer arithmetic function:

```

01. ; Demetrios Doumas Optimized code 11/6/15
02. .686P
03. .XMM
04. include listing.inc
05. .model flat
06.
07. _TEXT SEGMENT
08.
09. ;offset from base pointer
10. _p$ = -8 ; size = 4
11. _array$ = 8 ; size = 4
12. _size$ = 12 ; size = 4
13.
14. ?ClearUsingPointers@@YAXPAHH@Z PROC ; ClearUsingPointers, COMDAT
15.
16. ; Line 1
17. push ebp
18. mov ebp, esp
19. sub esp, 204 ; 000000ccH
20. push ebx
21. push esi
22. push edi
23. lea edi, DWORD PTR [ebp-204]
24. mov ecx, 51 ; 00000033H
25. mov eax, -858993460 ; ccccccccH
26. rep stosd
27. ;;;;;;;;;;;;;;
28. ; save the beginning address of A to _p$ or point _p to the address the first address of array
29. mov ecx, 0 ;
30. add ecx, DWORD PTR _array$[ebp]
31. mov DWORD PTR _p$[ebp], ecx
32.
33. ; need to save the last address to edx register
34. mov eax, DWORD PTR _size$[ebp] ; load the size to eax
35. mov ecx, DWORD PTR _array$[ebp] ; load the memory add of A into ecx
36. lea edx, DWORD PTR [ecx+eax*4]
37.
38. ; Condition
39. $loop:
40. cmp DWORD PTR _p$[ebp], edx
41. jae SHORT $exit

```

Continued...

```

27. ;;;;;;;;;;;;;;
28. ; save the beginning address of A to _p$ or point _p to the address the first address of array
29. mov ecx, 0 ;
30. add ecx, DWORD PTR _array$[ebp]
31. mov DWORD PTR _p$[ebp], ecx
32.
33. ; need to save the last address to edx register
34. mov eax, DWORD PTR _size$[ebp] ; load the size to eax
35. mov ecx, DWORD PTR _array$[ebp] ; load the memory add of A into ecx
36. lea edx, DWORD PTR [ecx+eax*4]
37.
38. ; Condition
39. $loop:
40. cmp DWORD PTR _p$[ebp], edx
41. jae SHORT $exit
42.
43. ; Body of the Loop
44. mov eax, DWORD PTR _p$[ebp] ; give the current address of _p put it into eax
45. mov DWORD PTR [eax], 0 ; initilaize it to zero
46. add eax, 4 ; Increment to the next value
47. mov DWORD PTR _p$[ebp], eax ; Move the new address to _p
48. jmp SHORT $loop
49.
50. $exit:
51. ;;;;;;;;;;;;;;
52.
53. ; Line 5
54. pop edi
55. pop esi
56. pop ebx
57. mov esp, ebp
58. pop ebp
59. ret 0
60. ?ClearUsingPointers@@YAXPAHH@Z ENDP ; ClearUsingPointers
61. _TEXT ENDS
62. END

```

The code in the last page and above is the optimized compiled generated assembly code. There are several of differences between the optimized compiler assembly code and the original compiler assembly listing code. In the complier assembly code there are 16 instruction and several of jumps or loops. The last address of the array is computed within the loop, where as the optimized code has the calculation of the last address outside the single loop. There are six instructions outside of the loop, where as in the complied generated code all the instruction has loops in between each other.

I followed the algorithm of the c++ code function and turn it to assemble code by editing the compile generated code. The first part is to declare the pointer. Make the pointer point to the first memory of address array. Next is to get the last address of the array. Then compare the current address with the last address of the array. If they are equal then jump to exit label and the stack begins to deallocate. If it does not exit, then increase the pointer and set the value of array at current index to zero. This program continues the same pattern until the current pointer points to the last address.

## Performance Measurements between the Use of Indexing and Pointer Arithmetic in Visual Studio:

Time Measurement code for measuring performance of the clearing array with Index

```
01. #include "stdafx.h";
02. #include <windows.h>;
03. using namespace System;
04. using namespace std;
05. #include <fstream>
06.
07. void ClearUsingIndex(int Array[], int size);
08.
09. static int A[100000000];
10.
11. int N[7] = {10, 100, 1000, 10000, 100000, 1000000, 10000000};
12.
13. int main()
14. {
15.
16.     int size = 10;
17.     _int64 ctrl = 0, ctr2 = 0, freq = 0;
18.     int acc = 0, i = 0;
19.
20.     double Avg = 0.0;
21.     ofstream myfile;
22.     myfile.open("IndexClearArraydata.txt");
23.
24.     for (int i = 0; i < 10; i++)
25.     {
26.         Console::WriteLine(A[i]);
27.     }
28.
```

continued...

```
28.
29.     double k = 0;
30.     for (int u = 0; u < 7; u++){
31.         for (int i = 0; i < 5; i++)
32.         {
33.
34.             if (QueryPerformanceCounter((LARGE_INTEGER*)&ctrl) != 0){
35.                 ClearUsingIndex(A, N[u]);
36.                 QueryPerformanceCounter((LARGE_INTEGER*)&ctr2);
37.                 QueryPerformanceFrequency((LARGE_INTEGER*)&freq);
38.
39.                 k = ("0", (((ctr2 - ctrl)* 1.0) / freq)) + k;
40.                 //Console::WriteLine("{0}", (((ctr2 - ctrl)* 1.0) / freq).ToString());
41.             }
42.
43.         }
44.         Avg = k / 5;
45.         myfile << Avg << "\t" << N[u] << "\n";
46.     }
47.
48.     Console::WriteLine(Avg);
49.
50.     for (int i = 0; i < 10; i++)
51.     {
52.         Console::WriteLine(A[i]); // << endl;
53.     }
54.     return 0;
55.
56.     myfile.close();
57.     Console::WriteLine();
58.
59.     Console::Read();
60.
61. }
```

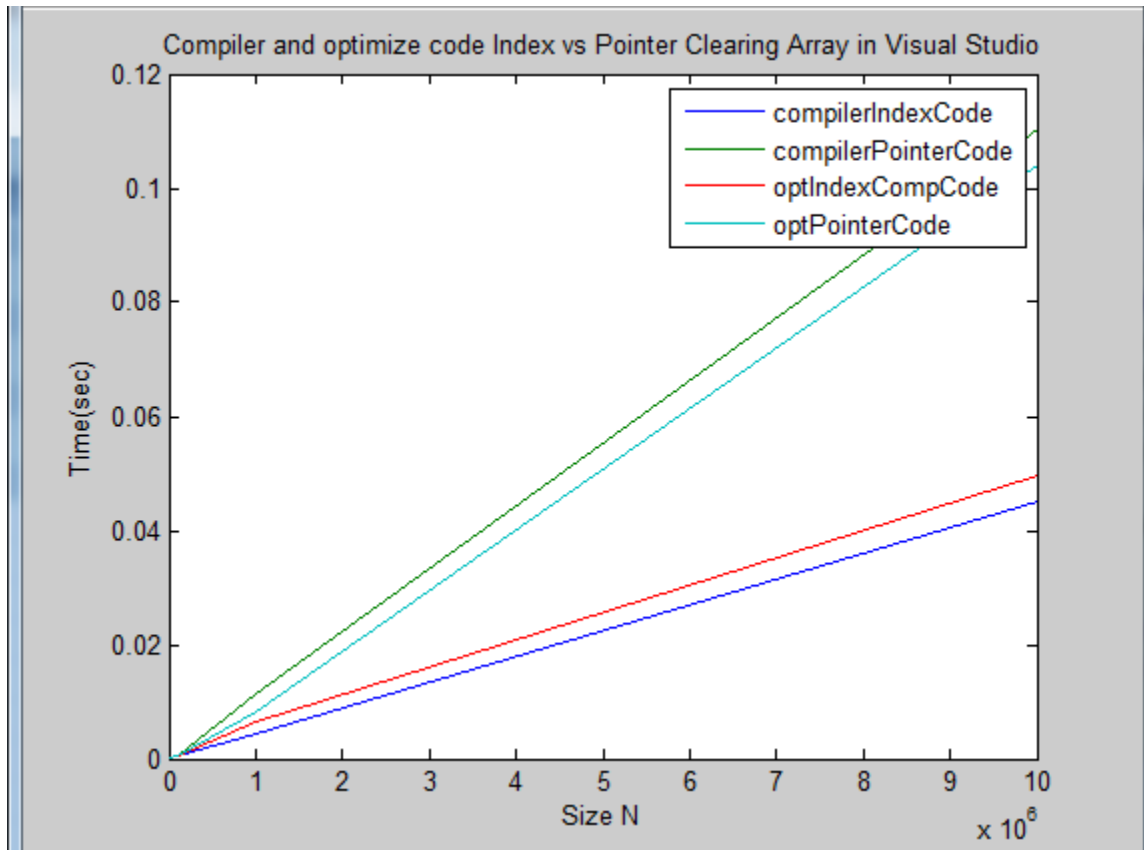
Time Measurement code for measuring performance of the clearing array with pointers:

```
01. #include "stdafx.h";
02. #include <windows.h>;
03. using namespace System;
04. using namespace std;
05. #include <fstream>
06.
07. void ClearUsingPointers(int *array, int size);
08.
09. static int A[100000000];
10.
11. int N[7] = { 10, 100, 1000, 10000, 100000, 1000000, 10000000 };
12.
13. int main()
14. {
15.
16.     int size = 10;
17.     _int64 ctr1 = 0, ctr2 = 0, freq = 0;
18.     int acc = 0, i = 0;
19.
20.     double Avg = 0.0;
21.     ofstream myfile;
22.     myfile.open("PointerClearArraydata.txt");
23.
24.     for (int i = 0; i < 10; i++)
25.     {
26.         Console::WriteLine(A[i]);
27.     }
28.
```

Continued...

```
29.     double k = 0;
30.     for (int u = 0; u < 7; u++){
31.         for (int i = 0; i < 5; i++)
32.         {
33.
34.             if (QueryPerformanceCounter((LARGE_INTEGER*)&ctr1) != 0){
35.                 ClearUsingPointers(A, N[u]);
36.                 QueryPerformanceCounter((LARGE_INTEGER*)&ctr2);
37.                 QueryPerformanceFrequency((LARGE_INTEGER*)&freq);
38.
39.                 k = ("0", (((ctr2 - ctr1)* 1.0) / freq)) + k;
40.                 //Console::WriteLine("0", (((ctr2 - ctr1)* 1.0) / freq).ToString());
41.             }
42.
43.         }
44.         Avg = k / 5;
45.         myfile << Avg << "\t" << N[u] << "\n";
46.     }
47.
48.     Console::WriteLine(Avg);
49.
50.     for (int i = 0; i < 10; i++)
51.     {
52.         Console::WriteLine(A[i]); // << endl;
53.     }
54.     return 0;
55.
56.     myfile.close();
57.     Console::WriteLine();
58.
59.     Console::Read();
60.
61.
62.
63.
64.
65. }
```

Total performance:



The results were unexpected. Both the compiler and optimized code using pointers should of have the fastest time to clear the array. However, as shown above the compiler generated code to clear the array using indexing was the fastest. The optimize pointer code to clear the array runs faster than the compiler generated code using pointers. The optimize code for clearing the array with indexing arithmetic ran slower than the complier generated code using indexing to clear the array.

The data is display in the table below:

Size N	Opt Time (clear by Pointer) sec	Opt Time (clear by Index) sec	Comp CodeTime (clear by Pointer) sec	Comp CodeTimeTime (clear by Index) sec
10	4.09E-05	3.39E-05	2.91E-05	3.56E-05
100	4.28E-05	3.56E-05	3.09E-05	3.76E-05
1000	4.90E-05	4.36E-05	3.79E-05	4.40E-05
10000	9.16E-05	8.67E-05	8.76E-05	8.23E-05
100000	0.000516884	0.000507771	0.000560606	0.000503266
1000000	0.0083963	0.00635917	0.0114575	0.00450646
10000000	0.104164	0.0496016	0.110457	0.044986

## Indexing VS Pointers in Linux Platform:

```
01. #include<iostream>
02. using namespace std;
03.
04. void ClearUsingIndex(int Array[], int size);
05.
06. static int Array[10] = {1,2,3,4,5,6,7,8,9,-1};
07.
08. int main() {
09.     int size = 10;
10.
11.
12.     for(int i=0;i<size;i++)
13.     {
14.         cout<<Array[i]<<endl;
15.     }
16.
17.     ClearUsingIndex( Array, size);
18.
19.
20.     for(int i=0;i<size;i++)
21.     {
22.         cout<<Array[i]<<endl;
23.     }
24.
25.     return 0;
26. }
```

```
01. void ClearUsingIndex(int Array[], int size) {
02.     int i;
03.     for (i = 0; i < size; i +=1)
04.         Array[i] = 0;
05. }
```

```
01. void ClearUsingPointers ( int *array, int size) {
02.     int *p;
03.     for (p = &array[0]; p < &array[size]; p = p + 1)
04.         *p = 0;
05. }
```

Linux Compiler Generated Assembly code below for index:

```
01. .file "indexcleararray.cpp"
02. .text
03. .globl _Z15ClearUsingIndexPii
04. .type _Z15ClearUsingIndexPii, @function
05. _Z15ClearUsingIndexPii:
06. .LFB0:
07. .cfi_startproc
08. pushq %rbp
09. .cfi_def_cfa_offset 16
10. .cfi_offset 6, -16
11. movq %rsp, %rbp
12. .cfi_def_cfa_register 6
13. movq %rdi, -24(%rbp)
14. movl %esi, -28(%rbp)
15. movl $0, -4(%rbp)
16. .L3:
17. movl -4(%rbp), %eax
18. cmpl -28(%rbp), %eax
19. jge .L4
20. movl -4(%rbp), %eax
21. cltq
22. leaq 0(,%rax,4), %rdx
23. movq -24(%rbp), %rax
24. addq %rdx, %rax
25. movl $0, (%rax)
26. addl $1, -4(%rbp)
27. jmp .L3
28. .L4:
29. nop
30. popq %rbp
31. .cfi_def_cfa 7, 8
32. ret
33. .cfi_endproc
34. .LFE0:
35. .size _Z15ClearUsingIndexPii, .-_Z15ClearUsingIndexPii
36. .ident "GCC: (Ubuntu 5.4.0-6ubuntu1~16.04.2) 5.4.0 20160609"
37. .section .note.GNU-stack,"",@progbits
```

The code above is the Linux compiler Generated assembly code listing. The stack frame is created on lines 7 to 12. The registers rdi and esi are stored in the stack with a certain offset from the base pointer. Lines 17 to 19 are the condition of the for-loop. The size of the array is stored in the base pointer at an offset of -28. The value of I gets stored in the register eax. The next three lines increments the address. After that the value in the address gets initialize to zero. The index value I get incremented. The pattern of the execution continues until the condition for loop when I is greater than the size and jumps to the few lines of code that starts to deallocate memory from the stack.

Linux Optimized compiler Generated Assembly code below for Index:

```
01. .text
02. .globl _Z15ClearUsingIndexPii
03. .type _Z15ClearUsingIndexPii, @function
04. _Z15ClearUsingIndexPii:
05. .LFB0:
06. .cfi_startproc
07. pushq %rbp
08. .cfi_def_cfa_offset 16
09. .cfi_offset 6, -16
10. movq %rsp, %rbp
11. .cfi_def_cfa_register 6
12.
13. movq %rdi, -24(%rbp)
14. movl %esi, -28(%rbp)
15. movl $0, -4(%rbp)
16.
17. .Loop:
18. movl -4(%rbp), %eax
19. cmpl -28(%rbp), %eax
20. jge .Exit
21. leaq 0(,%rax,4), %rdx
22. movq -24(%rbp), %rax
23. addq %rdx, %rax
24. movl $0, (%rax)
25. addl $1, -4(%rbp)
26. jmp .Loop
27.
28. .Exit:
29. nop
30. popq %rbp
31. .cfi_def_cfa 7, 8
32. ret
33. .cfi_endproc
34. .LFE0:
35. .size _Z15ClearUsingIndexPii, .-_Z15ClearUsingIndexPii
36. .ident "GCC: (Ubuntu 5.4.0-6ubuntu1~16.04.2) 5.4.0 20160609"
37. .section .note.GNU-stack,"",@progbits
```

The code above is the optimized version of the compiler generated assembly code for clearing the array using indexing arithmetic. Lines 6 to 12 is the creation of the stack frame. Lines 13 to 15 are different values stored on the stack at an offset from the base pointer. The index *I* is stored at an offset of -4 from the base pointer. Lines 18 to 20 are the condition of the for-loop. The register *eax* store *I* and is being compared to the array size of an offset of -28 from the base pointer. The rest of the code is the same. The only difference is that there is no *cltq* instruction and line 20 was remove because it was unnecessary.



Linux Compiler Generated Assembly code below with Pointers to clear array:

```
01. .file "pointercleararray.cpp"
02. .text
03. .globl _Z18ClearUsingPointersPii
04. .type _Z18ClearUsingPointersPii, @function
05. _Z18ClearUsingPointersPii:
06. .LFB0:
07. .cfi_startproc
08. pushq %rbp
09. .cfi_def_cfa_offset 16
10. .cfi_offset 6, -16
11. movq %rsp, %rbp
12. .cfi_def_cfa_register 6
13. movq %rdi, -24(%rbp)
14. movl %esi, -28(%rbp)
15. movq -24(%rbp), %rax
16. movq %rax, -8(%rbp)
17. .L3:
18. movl -28(%rbp), %eax
19. cltq
20. leaq 0(%rax,4), %rdx
21. movq -24(%rbp), %rax
22. addq %rdx, %rax
23. cmpq -8(%rbp), %rax
24. jbe .L4
25. movq -8(%rbp), %rax
26. movl $0, (%rax)
27. addq $4, -8(%rbp)
28. jmp .L3
29. .L4:
30. nop
31. popq %rbp
32. .cfi_def_cfa 7, 8
33. ret
34. .cfi_endproc
35. .LFE0:
36. .size _Z18ClearUsingPointersPii, .- _Z18ClearUsingPointersPii
37. .ident "GCC: (Ubuntu 5.4.0-6ubuntu1~16.04.2) 5.4.0 20160609"
38. .section .note.GNU-stack,"",@progbits
```

The code above is the compiler generated assembly code for clearing the array with pointer arithmetic. The stack frame is being created on lines 7 to 12. The values in line 13 to 16 are being stored on the stack by a particular offset from the base pointer. Lines 18 to 24 are the codes to increment the address of the pointer and compare the current address of the pointer to the last pointer of the array. Then the value zero gets copied to the current address which is stored in register rax. Then the execution of the code jumps back to incrementing the current address by 4 bytes. This pattern repeats until the register rax contains the last address of the array.

```

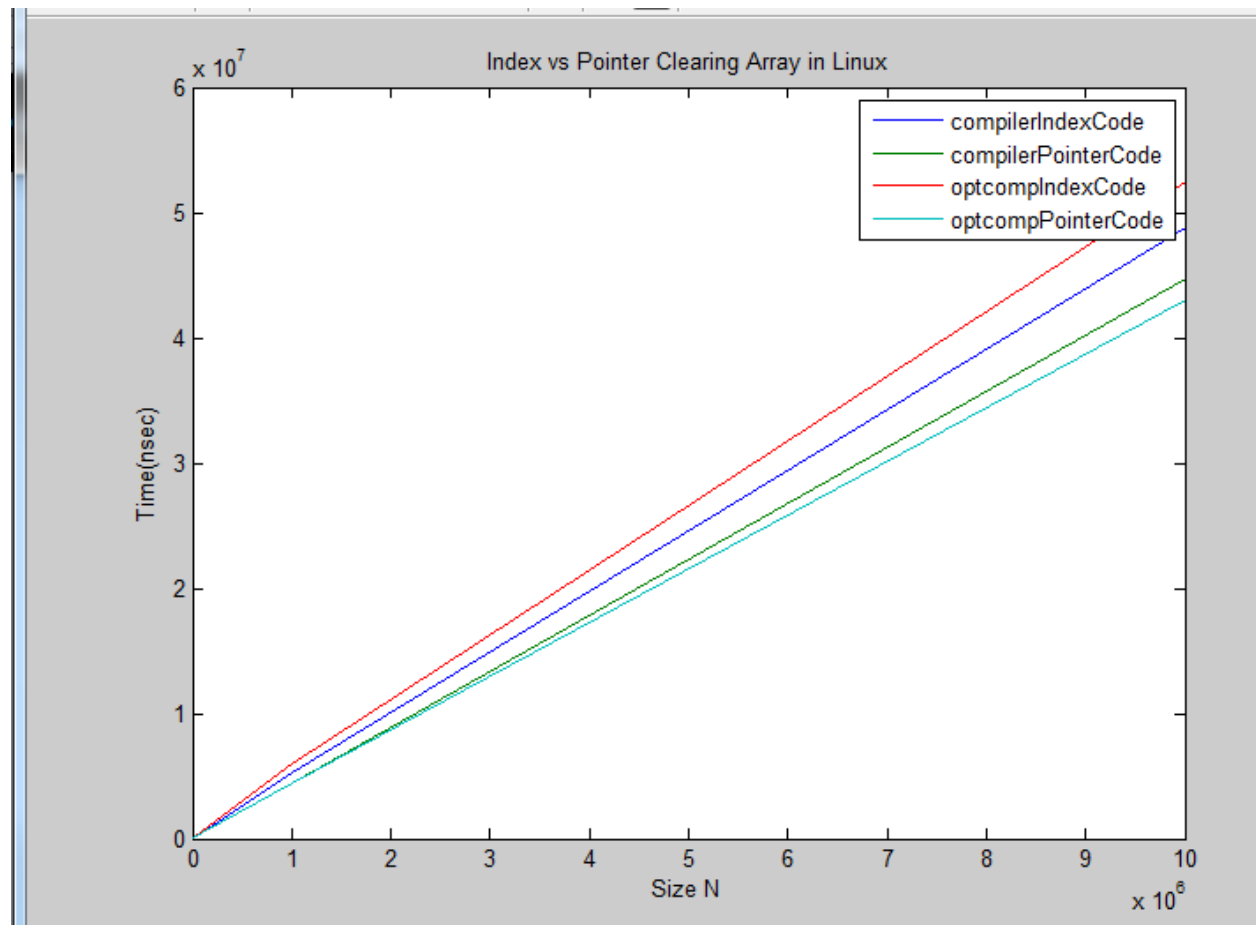
01.      .text
02.      .globl _Z18ClearUsingPointersPii
03.      .type _Z18ClearUsingPointersPii, @function
04.      _Z18ClearUsingPointersPii:
05.      .LFB0:
06.          .cfi_startproc
07.          pushq   %rbp
08.          .cfi_def_cfa_offset 16
09.          .cfi_offset 6, -16
10.          movq    %rsp, %rbp
11.          .cfi_def_cfa_register 6
12.
13.          movq    %rdi, -24(%rbp)
14.          movl    %esi, -28(%rbp)
15.          movq    -24(%rbp), %rax
16.          movq    %rax, -8(%rbp)
17.
18.      .Loop:
19.          movl    -28(%rbp), %eax
20.          leaq    0(,%rax,4), %rdx
21.          movq    -24(%rbp), %rax
22.          addq    %rdx, %rax
23.          cmpq    -8(%rbp), %rax
24.          jbe     .Exit
25.          movq    -8(%rbp), %rax
26.          movl    $0, (%rax)
27.          addq    $4, -8(%rbp)
28.          jmp     .Loop
29.
30.      .Exit:
31.          nop
32.          popq    %rbp
33.          .cfi_def_cfa 7, 8
34.          ret
35.          .cfi_endproc
36.      .LFE0:
37.          .size    _Z18ClearUsingPointersPii, . - _Z18ClearUsingPointersPii
38.          .ident   "GCC: (Ubuntu 5.4.0-6ubuntu1~16.04.2) 5.4.0 20160609"
39.          .section .note.GNU-stack,"",@progbits

```

Linux Optimized Compiler Generated Assembly code below with Pointers to clear array:

The code above is the optimized compiler generated assembly code for clearing an array with pointers. The code did not change much. Only the `cltq` instruction was removed. The rest of the code has the same functionality as the previous compiler generated code.

## Performance Measurements between the Use of Indexing and Pointer Arithmetic in Linux:



In the graph above shows the performance of the clearing the array functions using different techniques. The technique that perform the fastest is the clearing the array using pointers. The fastest performance was the optimized compiler generated assembly code using pointers. Clearing the array using indexing arithmetic took the longest to perform. The optimize code using indexing was slower than the compiler generated code using indexing arithmetic.

## Dot Product Computation Using Scalar code:

Main.cpp

```
01. #include<iostream>
02. #include<windows.h>;
03. using namespace std;
04.
05. int DotProduct(int x[], int y[], int xsize, int ysize);
06.
07. static int A[5] = { 1, 2, 3, 4, 5 };
08. static int B[5] = { 1, 2, 3, 4, 5 };
09. int main()
10. {
11.     int e = DotProduct(A, B, 5, 5);
12.     cout << e << endl;
13.
14.     system("Pause");
15.     return 0;
16. }
```

Dot Product using Indexing Technique:

```
01. int DotProduct(int x[], int y[], int xsize, int ysize)
02. {
03.     int result = 0;
04.     int j = 0;
05.     for (int i = 0; i < xsize; i++)
06.     {
07.         result = (x[i] * y[j]) + result;
08.         j++;
09.     }
10.
11.     return result;
12.
13. }
```

## Compiler Assembly listing code for Dot Product function using Index arithmetic:

```

01. ; Listing generated by Microsoft (R) Optimizing Compiler Version 18.00.21005.1
02.
03. .686P
04. .XMM
05. include listing.inc
06. .model flat
07.
08. INCLUDELIB MSVCRTD
09. INCLUDELIB OLDNAMES
10.
11. PUBLIC ?DotProduct@YAHQA0H0HH@Z ; DotProduct
12. EXTRN __RTC_InitBase:PROC
13. EXTRN __RTC_Shutdown:PROC
14. ; COMDAT rtc$TMZ
15. ;rtc$TMZ SEGMENT
16. ;__RTC_Shutdown.rtc$TMZ DD FLAT:__RTC_Shutdown
17. ;rtc$TMZ ENDS
18. ; COMDAT rtc$IMZ
19. ;rtc$IMZ SEGMENT
20. ;__RTC_InitBase.rtc$IMZ DD FLAT:__RTC_InitBase
21. ;rtc$IMZ ENDS
22. ; Function compile flags: /Odtp /RTCsu /ZI
23. ; COMDAT ?DotProduct@YAHQA0H0HH@Z
24. _TEXT SEGMENT
25. _i$1 = -32 ; size = 4
26. _j$ = -20 ; size = 4
27. _result$ = -8 ; size = 4
28. _x$ = 8 ; size = 4
29. _y$ = 12 ; size = 4
30. _xsize$ = 16 ; size = 4
31. _ysize$ = 20 ; size = 4
32. ?DotProduct@YAHQA0H0HH@Z PROC ; DotProduct, COMDAT
33. ; File c:\users\demetri\documents\visual studio 2013\projects\dotproduct\dotproduct\dotproduct.cpp
34. ; Line 2
35. push ebp
36. mov ebp, esp
37. sub esp, 228 ; 000000e4H
38. push ebx
39. push esi
40. push edi

```

Continued...

```

41. lea edi, DWORD PTR [ebp-228]
42. mov ecx, 57 ; 00000039H
43. mov eax, -858993460 ; ccccccccH
44. rep stosd
45. ; Line 3
46. mov DWORD PTR _result$[ebp], 0
47. ; Line 4
48. mov DWORD PTR _j$[ebp], 0
49. ; Line 5
50. mov DWORD PTR _i$1[ebp], 0
51. jmp SHORT $LN3@DotProduct
52. $LN2@DotProduct:
53. mov eax, DWORD PTR _i$1[ebp]
54. add eax, 1
55. mov DWORD PTR _i$1[ebp], eax
56. $LN3@DotProduct:
57. mov eax, DWORD PTR _i$1[ebp]
58. cmp eax, DWORD PTR _xsize$[ebp]
59. jge SHORT $LN1@DotProduct
60. ; Line 7
61. mov eax, DWORD PTR _i$1[ebp]
62. mov ecx, DWORD PTR _x$[ebp]
63. mov edx, DWORD PTR _j$[ebp]
64. mov esi, DWORD PTR _y$[ebp]
65. mov eax, DWORD PTR [ecx+eax*4]
66. imul eax, DWORD PTR [esi+edx*4]
67. add eax, DWORD PTR _result$[ebp]
68. mov DWORD PTR _result$[ebp], eax
69. ; Line 8
70. mov eax, DWORD PTR _j$[ebp]
71. add eax, 1
72. mov DWORD PTR _j$[ebp], eax
73. ; Line 9
74. jmp SHORT $LN2@DotProduct
75. $LN1@DotProduct:
76. ; Line 11
77. mov eax, DWORD PTR _result$[ebp]
78. ; Line 13
79. pop edi
80. pop esi
81. pop ebx

```

## Optimization of compiler Code using Index arithmetic for Dot Product computation:

```

01. .686P
02. .model flat
03. INCLUDELIB MSVCRTD
04. INCLUDELIB OLDNAMES
05. PUBLIC ?DotProduct@@YAHQAH0HH@Z ; DotProduct
06. _TEXT SEGMENT
07. _i$1 = -32 ; size = 4
08. _j$ = -20 ; size = 4
09. _result$ = -8 ; size = 4
10. _x$ = 8 ; size = 4
11. _y$ = 12 ; size = 4
12. _xsize$ = 16 ; size = 4
13. _ysize$ = 20 ; size = 4
14.
15. ?DotProduct@@YAHQAH0HH@Z PROC ; DotProduct, COMDAT
16. ; Line 2
17. push ebp
18. mov ebp, esp
19. sub esp, 228 ; 000000e4H
20. push ebx
21. push esi
22. push edi
23. lea edi, DWORD PTR [ebp-228]
24. mov ecx, 57 ; 00000039H
25. mov eax, -858993460 ; cccccccH
26.
27. rep stosd
28. mov DWORD PTR _result$[ebp], 0 ; result=0
29. mov DWORD PTR _j$[ebp], 0 ; j=0
30. mov DWORD PTR _i$1[ebp], 0 ; i=0
31.

```

Continue...

```

32. ;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;
33. ;for condition (i<size)
34. $Loop:
35. cmp eax, DWORD PTR _xsize$[ebp] ; compare index i with the size of the array
36. jge SHORT $Exit
37.
38. ;; increase i by one
39. mov eax, DWORD PTR _i$1[ebp] ; i is saved to eax
40. add eax, 1 ; eax increases i by 1
41. mov DWORD PTR _i$1[ebp], eax ; saves it back to i
42.
43. ;;;; result= X[i]*Y[i] + result;
44. mov eax, DWORD PTR _i$1[ebp] ; move i to eax
45. mov ecx, DWORD PTR _x$[ebp] ; move address of x to ecx
46. mov edx, DWORD PTR _j$[ebp] ; move j to edx
47. mov esi, DWORD PTR _y$[ebp] ; move address of y to esi
48. mov eax, DWORD PTR [ecx+edx*4] ; increment i and move it to eax
49. imul eax, DWORD PTR [esi+edx*4] ; y[j]
50. add eax, DWORD PTR _result$[ebp] ; add result to eax
51. mov DWORD PTR _result$[ebp], eax ; save eax back to results
52. ;;;; increase j
53. mov eax, DWORD PTR _j$[ebp] ; move j value to eax
54. add eax, 1 ; add 1 to eax
55. mov DWORD PTR _j$[ebp], eax ; save eax back to j
56. ;;; Save i to eax
57. mov eax, DWORD PTR _i$1[ebp] ; move i to eax
58. jmp SHORT $Loop
59.
60. $Exit
61. ;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;
62. ; Line 13
63. pop edi
64. pop esi
65. pop ebx
66. mov esp, ebp
67. pop ebp
68. ret 0
69. ?DotProduct@@YAHQAH0HH@Z ENDP ; DotProduct
70. _TEXT ENDS
71. END

```

In the last two pages there is a generated compiler assembly code and a corresponding optimized compiler generated assembly code for dot product computation using indexing arithmetic. The compiler generated code optimize code contains the variables i, j and result outside of the loop. The optimize code contains less number of loops that are in the compiler generated assembly code. They both follow the same functionality. Initialize the first three variables to zero. Check condition. The condition is whether or not the current index has reached to the size of the array. If it has not reach the end of the array, then the dot product computation continues to happen. If the index has reached the last index the function begins to exit and return the value result.

### Dot Product using Pointer Technique:

main.cpp

```
01. #include<iostream>
02. using namespace std;
03.
04. int DotProductPointers(const int* x,const int* y, int size);
05.
06. static int A[5] = { 1, 2, 3,4,5};
07. static int B[5] = { 1, 2, 3,4,5};
08. int main()
09. {
10.     int e = DotProductPointers(A, B, 5);
11.     cout << e << endl;
12.
13.     system("Pause");
14.     return 0;
15. }
```

dotproduct.cpp

```
01. int DotProductPointers(const int *a, const int *b, int size) {
02.
03.     const int *p, *q;
04.     int result = 0;
05.     for (p = a, q = b; p < a + size; p++, q++) {
06.         result += *p * *q;
07.     }
08.     return result;
09. }
```

## Compiler Assembly listing code for Dot Product function using Pointer arithmetic:

```

01. ; Listing generated by Microsoft (R) Optimizing Compiler Version 18.00.21005.1
02.
03. TITLE c:\Users\Demetri\documents\visual studio 2013\Projects\Project5\Project5\DotProductPointer.cpp
04. .686P
05. .XMM
06. include listing.inc
07. .model flat
08.
09. INCLUDELIB MSVCRTD
10. INCLUDELIB OLDNAMES
11.
12. PUBLIC ?DotProductPointers@@YAHPBH0H@Z ; DotProductPointers
13. EXTRN __RTC_InitBase:PROC
14. EXTRN __RTC_Shutdown:PROC
15. ; COMDAT rtc$TMZ
16. rtc$TMZ SEGMENT
17. ;__RTC_Shutdown.rtc$TMZ DD FLAT:__RTC_Shutdown
18. rtc$TMZ ENDS
19. ; COMDAT rtc$IMZ
20. rtc$IMZ SEGMENT
21. ;__RTC_InitBase.rtc$IMZ DD FLAT:__RTC_InitBase
22. rtc$IMZ ENDS
23. ; Function compile flags: /Odtp /RTCsu /ZI
24. ; COMDAT ?DotProductPointers@@YAHPBH0H@Z
25. _TEXT SEGMENT
26. _result$ = -32 ; size = 4
27. _q$ = -20 ; size = 4
28. _p$ = -8 ; size = 4
29. _a$ = 8 ; size = 4
30. _b$ = 12 ; size = 4
31. _size$ = 16 ; size = 4
32. ?DotProductPointers@@YAHPBH0H@Z PROC ; DotProductPointers, COMDAT
33. ; File c:\Users\demetri\documents\visual studio 2013\projects\project5\project5\dotproductpointer.cpp
34. ; Line 1
35. push ebp
36. mov ebp, esp
37. sub esp, 228 ; 000000e4H
38. push ebx
39. push esi
40. push edi
41. lea edi, DWORD PTR [ebp-228]
42. mov ecx, 57 ; 00000039H
43. mov eax, -858993460 ; ccccccccH
44. rep stosd
45. ; Line 4
46. mov DWORD PTR _result$(ebp), 0
47. ; Line 5
48. mov eax, DWORD PTR _a$(ebp)
49. mov DWORD PTR _p$(ebp), eax
50. mov ecx, DWORD PTR _b$(ebp)
51. mov DWORD PTR _q$(ebp), ecx
52. jmp SHORT $LN3@DotProduct
53. $LN2@DotProduct:
54. mov eax, DWORD PTR _p$(ebp)
55. add eax, 4
56. mov DWORD PTR _p$(ebp), eax
57. mov ecx, DWORD PTR _q$(ebp)
58. add ecx, 4
59. mov DWORD PTR _q$(ebp), ecx
60. $LN3@DotProduct:
61. mov eax, DWORD PTR _size$(ebp)
62. mov ecx, DWORD PTR _a$(ebp)
63. lea edx, DWORD PTR [ecx+eax*4]
64. cmp DWORD PTR _p$(ebp), edx
65. jae SHORT $LN1@DotProduct
66. ; Line 6
67. mov eax, DWORD PTR _p$(ebp)
68. mov ecx, DWORD PTR _q$(ebp)
69. mov edx, DWORD PTR [eax]
70. imul edx, DWORD PTR [ecx]
71. add edx, DWORD PTR _result$(ebp)
72. mov DWORD PTR _result$(ebp), edx
73. ; Line 7
74. jmp SHORT $LN2@DotProduct
75. $LN1@DotProduct:
76. ; Line 8
77. mov eax, DWORD PTR _result$(ebp)
78. ; Line 9
79. pop edi
80. pop esi
81. pop ebx
82. mov esp, ebp
83. pop ebp
84. ret 0
85. ?DotProductPointers@@YAHPBH0H@Z ENDP ; DotProductPointers
86. _TEXT ENDS
87. END

```



## Optimizing compiler Assembly listing code for Dot Product function using Pointer arithmetic:

```

01. .686P
02. .XMM
03. include listing.inc
04. .model flat
05.
06. INTEL LIB MSVCRTD
07. INTEL LIB OLDNAMES
08.
09. PUBLIC ?DotProductPointers@@YAHBPBH0H@Z ; DotProductPointers
10. EXTRN __RTC_InitBase:PROC
11. EXTRN __RTC_Shutdown:PROC
12.
13. _TEXT SEGMENT
14. _result$ = -32 ; size = 4
15. _q$ = -20 ; size = 4
16. _p$ = -8 ; size = 4
17. _a$ = 8 ; size = 4
18. _b$ = 12 ; size = 4
19. _size$ = 16 ; size = 4
20. ?DotProductPointers@@YAHBPBH0H@Z PROC ; DotProductPointers, COMDAT
21.
22. ; Line 1
23. push ebp
24. mov ebp, esp
25. sub esp, 228 ; 000000e4H
26. push ebx
27. push esi
28. push edi
29. lea edi, DWORD PTR [ebp-228]
30. mov ecx, 57 ; 00000039H
31. mov eax, -858993460 ; ccccccccH
32. rep stosd
33.
34. mov DWORD PTR _result$(ebp), 0 ; results =0
35. mov eax, DWORD PTR _a$(ebp) ;
36. mov DWORD PTR _p$(ebp), eax ; p points a
37. mov ecx, DWORD PTR _b$(ebp) ;
38. mov DWORD PTR _q$(ebp), ecx ; q points b
39.
40. ;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;
41. ; calculate the end of the array and compare current pointer to the end pointer
42.
43. $Loop:
44. mov eax, DWORD PTR _size$(ebp)
45. mov ecx, DWORD PTR _a$(ebp)
46. lea edx, DWORD PTR [ecx+eax*4]
47. cmp DWORD PTR _p$(ebp), edx
48. jae SHORT $Exit
49. ; Multiply result= (x[i] + y[i]) + result;
50. mov eax, DWORD PTR _p$(ebp)
51. mov ecx, DWORD PTR _q$(ebp)
52. mov edx, DWORD PTR [eax]
53. imul edx, DWORD PTR [ecx]
54. add edx, DWORD PTR _result$(ebp)
55. mov DWORD PTR _result$(ebp), edx
56. ; increment pointer p and q by 4 bytes
57. mov eax, DWORD PTR _p$(ebp)
58. add eax, 4
59. mov DWORD PTR _p$(ebp), eax
60. mov ecx, DWORD PTR _q$(ebp)
61. add ecx, 4
62. mov DWORD PTR _q$(ebp), ecx
63. jmp SHORT $Loop
64.
65. ;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;;
66. $Exit:
67. ; Line 8
68. mov eax, DWORD PTR _result$(ebp)
69. ; Line 9
70. pop edi
71. pop esi
72. pop ebx
73. mov esp, ebp
74. pop ebp
75. ret 0
76. ?DotProductPointers@@YAHBPBH0H@Z ENDP ; DotProductPointers
77. _TEXT ENDS
78. END

```

In the last two pages is the implementation of the dot product using pointer arithmetic. On page 24, is the compiler generated assembly code. On page 25, is the optimized compiler generated assembly code. Both codes are similar to each other. There were no other ways to optimize the code. There are less jumps in the optimize code then the compiler generated code. They have the same functionality. They both have pointer p and q point to a and b respectively. Then the size of the array is calculated. The size is needed to compute the end of the array. The memory location of the end is being compared to the current pointer in p. If the pointer p has not reach the end, then the dot product computation is continued.

## Dot Product Computation Using Vector Instruction:

```
1. // VectorInstructionDotProduct.cpp : main project file.
2. #include "stdafx.h"
3. extern "C" {
4. #include<xmmintrin.h>
5. }
6.
7. #include<iostream>
8. #include<fstream>
9. #include<windows.h>
10.
11. using namespace std;
12. using namespace System;
13.
14. void DotProduct(float x[], float y[], int size);
15.
16. _declspec(align(16)) static float A[10000000];
17. _declspec(align(16)) static float B[10000000];
18. // 8 25 250 2500
19. int N[7] = { 8, 100, 1000, 10000, 100000, 1000000, 10000000 };
20.
21. int main()
22. {
23.     int size = 10;
24.     _int64 ctrl = 0, ctr2 = 0, freq = 0;
25.     int acc = 0, i = 0;
26.
27.     double Avg = 0.0;
28.     ofstream myfile;
29.     myfile.open("VectorDotProductdata.txt");
30.
31.
32.     double k = 0;
33.     for (int p = 0; p < 10000000; p++)
34.     {
35.         A[p] = 1;
36.         B[p] = 1;
37.     }
38.
39.
40.     for (int u = 0; u < 7; u++){
41.         for (int i = 0; i < 5; i++)
42.         {
43.
44.             if (QueryPerformanceCounter((LARGE_INTEGER*)&ctrl) != 0){
45.                 DotProduct(A, B, N[u]);
46.                 QueryPerformanceCounter((LARGE_INTEGER*)&ctr2);
47.                 QueryPerformanceFrequency((LARGE_INTEGER*)&freq);
48.
49.                 k = ("0", (((ctr2 - ctrl)* 1.0) / freq)) + k;
50.             }
51.
52.         }
53.         Avg = k / 5;
54.         myfile << Avg << "\t" << N[u] << "\n";
55.     }
```

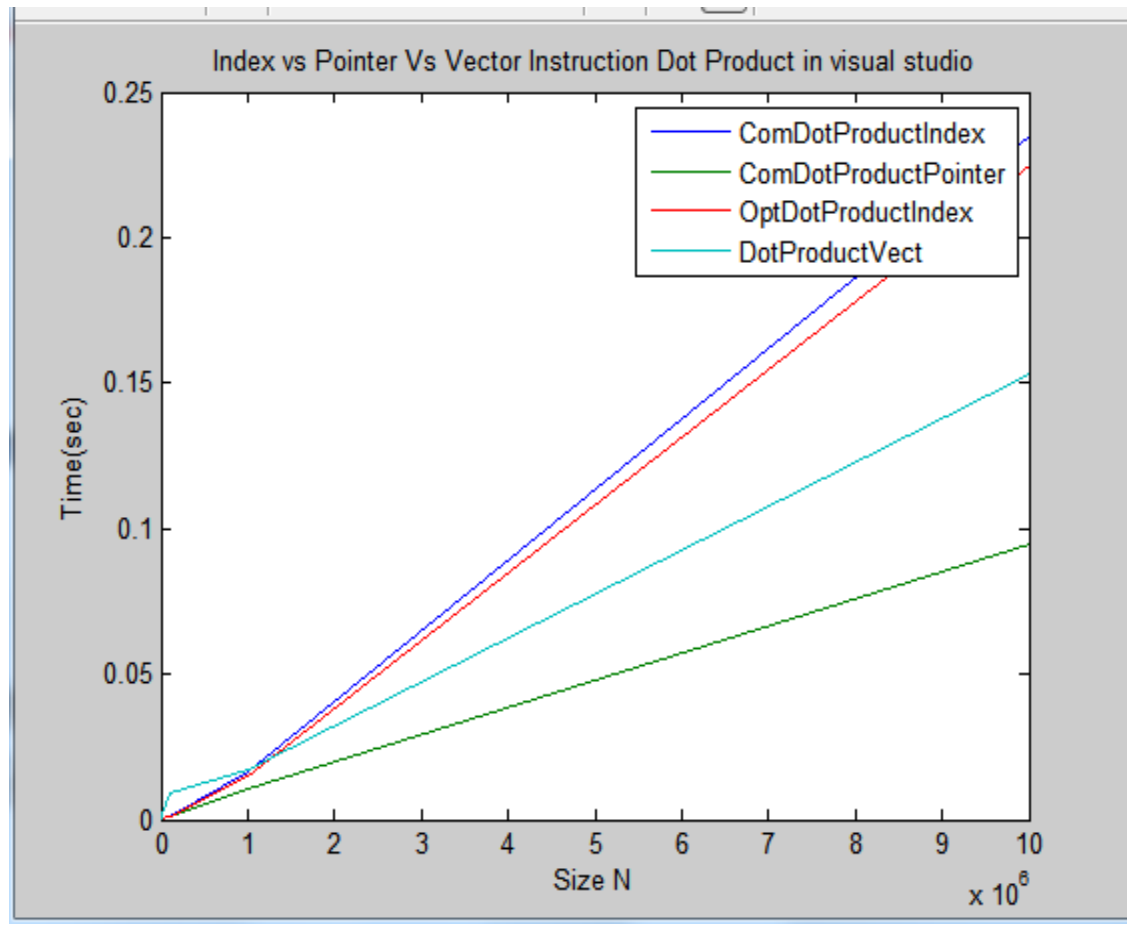
```

56.     }
57.
58.     Console::WriteLine(Avg);
59.     myfile.close();
60.     Console::WriteLine();
61.
62.     Console::Read();
63.
64.
65.     return 0;
66. }
67.
68.
69. void DotProduct(float x[], float y[], int size)
70. {
71.     __m128 temp1, temp2, res;
72.     float sum = 0.0;
73.     _declspec(align(16)) float temp4[4];
74.
75.     for (int e = 0; e < size; e += 4)
76.     {
77.         temp1 = _mm_load_ps(x + e);
78.         temp2 = _mm_load_ps(y + e);
79.         res = _mm_mul_ps(temp1, temp2);
80.         res = _mm_hadd_ps(res, temp2);
81.         res = _mm_hadd_ps(res, temp2);
82.         _mm_store_ps(temp4, res);
83.         sum += temp4[0];
84.
85.     }
86.
87.     cout << sum << endl;
88.
89. }

```

In the code above is the implementation of dot product using vector instructions. The inputs to the dot product functions are two arrays. The two arrays are turned into vector of size N. The loop reads every four values of each vector until size N in groups of four. The `_mm_mul_ps` command multiplies each group of four from both vectors and inputs them into result (res). The result is added horizontally twice to perform the sum of the dot product. The next instruction `_mm_store_ps` stores the value of result to a temporary array of four. The result of the dot-product of every four groups is stored at the first index of temp four. Then variable sum acts like an accumulator and adds all the result from each group of four to produce a single sum. The code above iterates different values of N (the size of the two arrays). It calculates the time five times and computes the avg. The data is then stored to an external text file.

## Performance Measurements:



The performances of the dot product with different implementation are shown above. The fastest implementation of the dot product is the compiler generated assembly code using pointer arithmetic. The second fastest is the dot-product computation using vector instructions. Ideally the vector instruction implementation should be the fastest in terms of performance. One reason that it is not the fastest because the loop to multiple and adds in groups of four makes it sequential statement and not in parallel with the other elements of the array. This is not the ideal parallel performance that is expected with vector instructions. The loop makes the calculation sequential in parallel groups of four instead of size multiply groups of four computations in parallel.

## Conclusion:

The compiler does not always generate the best assembly code. The fastest way to clear an array by initialize each element to zero and dot product computation is using pointer arithmetic. The second fastest performance for calculating the dot product is the implementation of vector instructions. The compiler in Linux operating system appears to generated assembly code better than the visual studio compiler. However, the assembly code generated by the visual studio compiler generates assembly code that it is easier to understand.

I learned a lot with this take home test. I learned to generated assembly compiler code and run it side by side with a c++ file for both windows and Linux operating systems. I learn that you can optimize compiler generated code. I also learned about vector instructions and how to implement it. The most importantly I learned to calculate time in visual studio by query-performance counter and in Linux get-time.

# Appendix:

Generated Measurement plots (Matlab):

```
1. %% Take home test 3
2. %% Plot Index vs Pointer Clearing Array in visual studio
3. x=[10, 100, 1000, 10000, 100000, 1000000, 10000000];% Size N
4. compilerIndexCode=[3.56331e-005,3.75785e-005,4.40294e-005,8.23247e-
005,0.000503266,0.00450646,0.044986 ];% time seconds
5. compilerPointerCode=[0.0000290799,0.0000309229,0.0000378857, 0.0000876491,0.000560606,0
.0114575,0.110457];
6. optIndexCompCode=[3.39E-05,3.56E-05,4.36E-05,8.67E-
05,0.000507771,0.00635917,0.0496016];
7. optPointerCode=[4.09E-05,4.28E-05,4.90E-05,9.16E-05,0.000516884,0.0083963,0.104164];
8.
9. plot(x,compilerIndexCode,x,compilerPointerCode,x,optIndexCompCode,x,optPointerCode)
10. xlabel('Size N')
11. ylabel('Time(sec)')
12. legend('compilerIndexCode','compilerPointerCode','optIndexCompCode','optPointerCode')

13. title('Compiler and optimize code Index vs Pointer Clearing Array in Visual Studio')
14.
15.
16. %% Plot Index vs Pointer Clearing Array in Linux
17. x=[10, 100, 1000, 10000, 100000, 1000000, 10000000];% Size N
18. compilerIndexCode=[16020.2,3157,9219.4,50928.4, 461470,5.21314e+06, 4.88297e+07];% time
seconds
19. compilerPointerCode=[1242.8, 2877.2, 9707.6,44502.8,427304, 4.49783e+06, 4.47219e+07];%
time seconds
20. optcompIndexCode=[1257.2, 3324.8,12711.6,87693.4,672992,6.00486e+06, 5.2492e+07];% time
seconds
21. optcompPointerCode=[1326.8,2793.2,7961,103061,535444,4.43619e+06,4.30305e+07];% time se
conds
22. plot(x,compilerIndexCode,x,compilerPointerCode,x,optcompIndexCode,x,optcompPointerCode)
;
23. legend('compilerIndexCode','compilerPointerCode','optcompIndexCode','optcompPointerCode
')
24. xlabel('Size N')
25. ylabel('Time(nsec)')
26. title('Index vs Pointer Clearing Array in Linux')
27. %% Plot Index VS Pointer Scaler and vector Dot Product in Visual Studio
28. x=[10, 100, 1000, 10000, 100000, 1000000, 10000000];% Size N
29. ComDotProductIndex=[4.9661e-005,5.21185e-005,6.28698e-
005,.000148778, .00106705,.0167145,.234933];% time seconds
30. ComDotProductPointer=[3.82953e-005,4.0548e-005,4.97634e-
005,0.000141713,0.00091366,0.0109088,0.0948066];
31. OptDotProductIndex=[3.6657e-005,3.87049e-005,4.77155e-
005,0.000116524,0.00079857,0.0150687,0.225133];
32. DotProductVect=[0.000529379,0.00091817,0.00128884,0.00175658,0.00911156,0.0168489,0.153
381];
33. plot(x,ComDotProductIndex,x,ComDotProductPointer,x,OptDotProductIndex,x,DotProductVect)
;
34. %axis([0 100000000 0 1000000000000]);
35. xlabel('Size N')
36. ylabel('Time(sec)')
37. legend('ComDotProductIndex','ComDotProductPointer','OptDotProductIndex','DotProductVect
')
38. title('Index vs Pointer Vs Vector Instruction Dot Product in visual studio')
```